

# Part 1 report:

Tamar Michelson 323805861

Shilo Avital 206487407

The parameters of the best model for both tasks (NER and POS):

Hidden dim = 100

Batch size = 32

Num epochs = 10

Learning rate =  $1e-3$


Weight decay =  $1e-5$

## Considerations:

Words that appear in the dev sets but not in the training set are assigned the embedding for <UNK>, a special unknown token.

At the beginning and end of sentences (to build the sliding window), we pad the sequence with <PAD> tokens. These have a fixed embedding of zeros.

The ner plots:



The pos plots:

